

Daniel Pascal Lamblin

Brooklyn NY U.S.A. & Remote

m:dlamblin@gmail.com

t:+1 646 644 3587

gh:dlamblin.github.io

in:linkedin.com/in/daniellamblin

Skills *JVM:* Java, Guava, Guice, Dagger, SpringBoot, Lucene, GWT, Scala & Kotlin; *Unix:* Python, Go, C, Perl & Bash; *C#:* ASP.Net; *Data:* Airflow, MySQL, Hive, Presto, Spark, Megastore, Beam/Cloud Dataflow; *Web:* NodeJS, Vue.js, HTML, CSS & JavaScript; *AWS:* ALB, ASG, EBS, EC2, EFS, EKS, EMR, Lambda, RDS, CloudTrail, CloudWatch, CloudFormation, S3 ...; Agile & TDD

Education Andrew Ng's Stanford ML-Class.org; **Online** – Became Coursera 2011–2011
Columbia University; **New York, NY** – Non-Degree Semester 2007–2007
Worcester Polytechnic Institute; **Worcester, MA** – BS Computer Science 1995–1999

Experience **Coupage**, *Senior Backend Software Engineer* Clickstream Web Logging **Seattle** 2020–Present

(Selected) ▶ Developed Customer Experience Analytics Platform tool for self-service funnel, journey and trend analysis of web logging data
▶ MVP for CXAP tool: Stood up the web front on Vue.js 2 for customers of Business Analytics team
▶ Designed the scheduling of CXAP jobs
▶ Delivered Journey and Trend analysis for backend of CXAP with Spark and Scala
▶ Mentored intern project to port and extend UI on Vue.js 3 & converted to hire
▶ Assisted Retail Delivery team's mission-critical multi-AZ expansion projects

Senior Data Engineer Data Platform Tools and Infra **Seattle** 2019–2020

- ▶ Migrated enterprise data warehouse's ~4000 Airflow ETL jobs from Airflow 1.8 to 1.10+, with custom semi-automated tooling.
- ▶ Right-size teams' infrastructure resource usage by monitoring for efficient use of EMR and Spark jobs
- ▶ Split Airflow into multiple deployments and onboard users to latest versions
- ▶ On call rotation as in Seoul.

Big Data Engineer Big Data Platform **Seoul** 2017–2019

- ▶ Onboard teams (edw, growth, pricing, catalog, search, retail etc.) into airflow; configure and maintain EMR clusters for teams' jobs
- ▶ Develop Data Platform Portal tool with Cluster start-stop API, management, and Data Discovery
- ▶ Scale out Airflow to ~6000 dags with multiple deployments
- ▶ On call rotation for Airflow, Presto, Hive, Hue, Zeppelin, Spark, HDFS, Zookeeper etc.

Data Engineer Data Platform **Seattle** 2016–2017

- ▶ Cloud migration. Lift and Shift on-prem IBM Netezza and HDFS Hive to EMR Hive on S3
- ▶ Update ETL into Airflow from Oozie and Talend
- ▶ Monitor data readiness with oncall rotations

Insight Data Science, *Data Engineering Fellow* **New York City** 2015–2016

- ▶ Realtime and batch processing on NYC MTA's GTFS stream with Kafka, Spark, HDFS, HBase, S3
- ▶ Generated user data in large scale for testing aim of notifying users of train delays
- ▶ Cluster on AWS EC2; project information and presentation at dlamblin.github.io/mta-delay-monitoring

Google, *Software Engineer* **New York City** 2010–2014

- ▶ Ported legacy Studio product, a rich advertisement authoring and QA web-app, to Google's Web Toolkit front-end with a stubby rpc backend and megastore datastore
- ▶ Developed a dashboard to track component usage data of ads by comparing html5 vs. flash authoring, common formats, layouts, and generated impressions. Utilized cross-team apis and internal versions of GFS, Cloud Dataflow and Drill
- ▶ Implemented critical preview features for monitoring ad unit interactions and compliance
- ▶ Migrated user records and assets to support new multi-account users and unified asset library view
- ▶ Reduced reprocessing and conserved storage of assets by fingerprinting uploads, both on individual files and within archives

Travelocity, *Senior Web Application Developer*

New York City 2006–2009

- Established features for IgoUgo.com using the ASP.NET 2.0 framework, C#, PrototypeJS
- Introduced a Lucene based index of content to offload db search as a wsdl service in Java and Spring with auto-completed suggestions for key geo-entities
- Boosted traffic ten fold through optimization of page structure, URLs and image file names

EMC Corp, *Associate Software Engineer*

Milford, MA 2000–2001

- Extended EMC Data Manager Volume & Tape Library Manager's media duplication processes targeting Petabyte capable systems like Sony PetaSite with multiple robots and drives
- Improved and maintained EDM as a multi-process C based system with Sun RPC, threading, and IPC signals & pipes; Resolved deadlock by refactoring mutex hierarchy